

SSHVisual: Visual Tool to Analyze SSH Data

Shreeman Gautam*, Robert Ricci*

*University of Utah

Abstract

Network operators use visualizations to detect attacks on their systems, exploring and interacting with data to uncover patterns that would otherwise stay hidden. Operators then use these findings to block attackers and to expand the knowledge base of known threats. Detecting anomalies and uncovering deeper patterns iteratively in data is at best, tedious and at worst, not achievable without data visualizations. Previous works in network visualization analyzed the IP and AS dimensions of SSH logins without much focus on the username aspect of the attackers' attempts, detected anomalous network events only where baseline events were available, and enabled operators to share findings with one another. Little attention has been paid to whether attackers observed at different times represent the same persistent actor. We present an iterative visualization SSHVisual, using login attempts collected on our own SSH-facing systems. SSHVisual analyzes the IP, AS, and username dimensions of SSH attacks, in service of SSH campaign detection. It complements automated detection by allowing operators to find anomalous events without normal baseline events. It correlates SSH attacks with real-world events, and it introduces a similarity metric that quantifies attacker persistence over time. Like previous systems, SSHVisual supports iterative and collaborative sharing of findings between operators. We evaluated SSHVisual using case studies and found that it helps us block attacks and uncover patterns. We identified campaigns in which a set of attempted usernames recurred across unrelated IPs and ASes. The similarity metric surfaced attackers returning across different time periods despite changing network origin. SSHVisual shows that treating usernames as a first-class analytic dimension reveals attacker coordination that IP-centric analysis misses, and that exploratory, operator-driven detection remains valuable.

I Introduction

Computer systems worldwide employ SSH [1] to allow their users to securely connect to them in order to execute commands and do remote administration. A lot of users do a username, password combination (password based authentication) in order to gain entry to the SSH system. Attackers, who are adversarial in nature, can use brute force attack mechanisms [2] to get into SSH systems. Sometimes, they guess what the username, password combinations are and other times, they take advantage of poorly secured systems to build botnets [2] to carry out further attacks [3, 4] or steal resources from the

systems, like cryptomining [5, 6].

Given the prevalence of SSH attacks in the IPv4 space [2], we collect SSH login data on our own systems and present SSHVisual. SSHVisual serves two kinds of network operators. Network administrators, who handle day-to-day SSH traffic, need to identify attackers quickly and block them. Security analysts, who study traffic over longer time periods, look for patterns that are not immediately apparent and share what they find across organizations. Consider a single time period: an administrator can isolate the highest-volume attacker and block it immediately, while an analyst examining the same window can iteratively surface a set of low-frequency attackers that attempt an identical sequence of usernames, and group them as a coordinated botnet.

Surveys of network visualization [7, 8, 9, 10, 11] establish the maturity of the field and the gaps we target. Our contributions follow from treating the username as an attacker identity. First, actor-level analysis: we treat the attacker as a first-class actor — including the username identity, not only the source IP or autonomous system (AS) — and we identify coordinated campaigns. Since network data affords no normal baseline for automated detection [12], this identification is necessarily exploratory (therefore, not antagonistic but complementary to automated detection with normal baselines) — operators recognize coordination visually rather than specifying it in advance. Second, a similarity metric that quantifies attacker similarity and lets operators track how actors evolve in the username dimension across time periods. Third, correlation of actor behavior with real-world events, which is more tractable through usernames than IPs, given the limitations of IP corroboration. Additionally, following established practices, we allow for iterative operator driven collaboration and labeling of irregularities found in the data.

The rest of the paper is as follows: Section II is Related Work, Section III is Tasks and Requirements, Section IV is Design, Section V is Evaluation and Section VI is Conclusion and Future Work.

II Related Work

A long line of tools [13, 14, 15, 16, 17, 18, 19, 20] render network traffic and security events navigable for operators and analysts with the majority being network flow and traffic oriented. Additionally, a smaller alert- and log-oriented group [21, 22, 23, 24, 25, 26, 27, 28, 29, 30] organizes IDS output and heterogeneous logs of which [23] contributes a visual analysis (of the CICIDS2017 dataset) that isolates at-

tack clusters — including SSH brute-force (SSH-Patator) — through clustering and dimensionality reduction, namely a multi-resolution wavelet-feature analysis.

A Attack Campaigns with Time

Honeypot literature does not promote the username identity as an actor in its own right (for attack campaign detection purposes): it remains an attribute to count, never a unit to track — the unit that, alongside the IP and AS, anchors our analysis. Some papers [31, 32, 33, 34, 35, 36, 37] collect and chart honeypot traffic, with several of them capturing SSH credentials (usernames and passwords). Where credentials are charted — the SCADA SSH honeypot [34], kippo-graph [35], and the container deployment [36] — usernames and passwords appear only as static top-N tallies (kippo-graph’s top-10 username and password bars are the archetype), never placed on a time axis even where the same tool charts attack volume over time. The WAN honeypot framework [38] is the closest prior system on the credential axis: it computes an IP-to-IP dictionary-overlap (containment) score to probe whether six same-/24 source IPs share a wordlist — a manual, one-off prototype of credential-based coordination analysis, performed on the IP unit and without a temporal dimension.

B Attacker Similarity with Time

Computed similarity between entities is well established in literature [39, 40, 41, 42]. We do not claim similarity computation itself as novel; the contribution lies in the entity it operates on and how it moves through time. On the entity: these systems cluster internal hosts [39], behavioral profiles [41], IP/port members [42], or internal enterprise users grouped by shared resources — not external attacker identities [40]. None treats the username as an attacker actor. On time: [41, 42] reach into the temporal dimension, but both do not compute the similarity of the attacker population across consecutive windows. We apply similarity based metrics — for example, Jaccard overlap of attackers between successive windows — particularly to the username actor unit, capturing how the attacker population persists or churns over time.

C Limits of automated detection

One premise of our approach is that SSH honeypot traffic does not admit the normal baseline that automated anomaly and campaign detection presuppose. The broader literature already documents the fragility of such detection: Sommer and Paxson [43] note that machine-learning anomaly detection is rarely deployed for intrusion detection given the variability of “normal” the cost of errors, and the difficulty of sound evaluation, while Axelsson [44] shows it is throttled by the base-rate fallacy, where the false-positive rate becomes the limiting factor once intrusions are rare relative to benign traffic. Our setting departs from that one entirely, in the direction Gates and Taylor [12] anticipate when they question whether attacks are rare or a stable normal exists, note that normal traffic may itself

be malicious and malicious traffic unremarkable, and observe that detection ultimately devolves to an operator investigating clusters of activity by hand. The consequence is visible in the systems cited above: Situ [39] learns a probabilistic model of each IP’s behavior and CLIQUE [45] scores each actor against a baseline of its own expected activity, so both presume a per-actor “normal” that near-uniformly hostile, constantly churning traffic cannot supply. Situ nonetheless concedes that automation alone is insufficient — operators “do not want another automated tool with algorithms they do not trust” — and builds a human-in-the-loop system; we reach the same conclusion from a starker starting point, lacking even the per-actor baseline it assumes, and so treat campaign identification as operator-driven from the outset.

D Actor identity and real-world event correlation

The attacks we observe do not arrive at a steady rate; they rise, fall, and change in character as events unfold in the wider world, so understanding them means lining them up against what is happening outside the network. The link is well documented: Bilge and Dumitras [46] found that the number of attacks against a given vulnerability can jump by up to five orders of magnitude — a hundred-thousand-fold increase — once it is publicly disclosed, and Sabottke et al. [47] showed that future exploitation can be predicted from signals in disclosure channels alone. Allodi and Massacci [48] add that a vulnerability’s official severity score is a poor guide to whether it is actually attacked, so the dependable way to learn what attackers are reacting to is to watch real activity and line it up against real events, not to read it off a score. We want to do this at the level of a campaign — a group of attackers acting together — rather than one attacker at a time, and that raises a practical problem: a campaign rarely shows up as a single source address. IP addresses do not stay tied to one machine — one host is assigned many addresses over time [49], hosts cycle through new addresses [50], and counting by address can overstate the number of real machines by orders of magnitude [51] — so it is likely that a view organized by address scatters one campaign across dozens of entries that look unrelated. The username pulls those pieces back together: the same login names likely recur across the changing addresses, so grouping by username re-assembles the campaign that an address-by-address view splits apart. A username is not a perfect identity either — names like “root” are tried by almost everyone, and one campaign runs through a whole list of names — so we treat it as one analysis dimension alongside the address and the network it belongs to (its AS), resolving the overlaps with the similarity and grouping methods above.

E Operator labeling and collaboration

Since detection on our data is operator-driven, the findings an operator produces by hand must be captured and shared rather than discarded. In two lines of work [52, 19], interactive labeling has been used to fold the operator’s judgments back into

the system. Collaborative analysis has been supported through shared artifacts with [17] and [20] being notable examples. These systems establish that labeling and collaboration are valuable, which is why our labeling-and-sharing mechanism is the step that turns an identified set of usernames into a durable, shared campaign.

III Tasks and Requirements

The design of SSHVisual draws on the authors’ experience as network operators and on gaps we identified in prior work on SSH attack analysis and security visualization. From these we derive three tasks and six requirements. Table 1 summarizes which requirements support which tasks.

A Tasks

T1: Identify SSH campaigns. Security analysts need to find groups of attackers that behave alike across usernames, IP addresses, autonomous systems (ASes), and countries, and to record those groups so they can be referenced later.

T2: Detect anomalies that warrant action. Network administrators need to notice unusual activity, such as a sudden spike or a single source dominating traffic, quickly enough to decide whether to block it.

T3: Share findings across operators. Analysts and administrators on the same network need to see what their colleagues have already investigated, so that the same patterns are not identified twice.

B Requirements

R1: Attacker similarity profiling. The system should be able to algorithmically cluster similar attackers based primarily along the username dimension, secondarily along IP and AS information and tertiarily along other dimensions. R1 mostly serves T1 as it allows analysts to build a repository of similar attackers.

R2: Sorting and ranking SSH attack data is heavy-tailed: a few usernames [53] and IP addresses [54] account for most attempts, and the remaining entities are buried beneath them. The system should therefore list every entity in each dimension, not only the most active ones, and let the operator sort and rank entities by any combination of attributes. R2 supports T1 and T2: administrators can surface anomalous entities that never reach the top of a volume ranking, and analysts can find entities with similar attribute values.

R3: Visual clarity. Line charts become harder to read accurately as more series are added [55]. Each linechart should therefore show at most ten entities. An aggregate chart of total attacks should accompany the per-entity charts, so that operators can judge how much of the total any one entity accounts for. For example, if 150M attacks occur on a given day and 75M of them come from a single username, that username should stand out when the two charts are compared. R3 primarily supports T2.

Table 1: Mapping of requirements to tasks (● primary support, ○ secondary support).

	R1	R2	R3	R4	R5	R6
T1: Identify SSH campaigns	●	●		○	●	●
T2: Detect anomalies		●	●		●	●
T3: Share findings				●		

R4: Shared tagging. Operators should be able to tag attackers, record both what a tag means and why it was applied to a particular attacker, and have those tags persist and remain visible to every operator using the system. Each tag application should record who applied it and when, so that later operators can judge whether a finding is still current. R4 primarily supports T3, and it supports T1 because tags accumulate into a repository of related attackers.

R5: Complementary tabular and temporal views. No single visualization serves every need. A table supports precise lookup and ranking but discards temporal structure; a linechart preserves temporal structure, showing spikes and campaigns at a glance, but conveys values less precisely. The system should provide both views and let operators move between them. An administrator may find and act on an anomaly using the table alone, whereas an analyst may use the table to assemble a set of related attackers and then examine their behavior over time in the linecharts. R5 supports T1 and T2.

R6: Interactive exploration of the linecharts. The linecharts should follow Shneiderman’s Visual Information-Seeking Mantra: overview first, zoom and filter, then details-on-demand [56]. They should open with an overview of the full time range, let the operator zoom into a time window and filter which entities are displayed, and reveal details about an individual entity on request. R6 supports T1 and T2: zooming isolates the time window of a suspected campaign, and filtering removes dominant entities that would otherwise mask smaller ones.

IV Design

SSHVisual combines two visual components, a discovery table and a set of linecharts, with two non-visual components, similarity search and shared tagging. We describe the visual design in Section A and the non-visual design in Section B. All figures show SSH login attempts collected between 1 November 2022 and 8 January 2023, a period of 69 days.

A Visual Design

A.1 Discovery Table

The interface provides two views of the data. The discovery table (Figure 1) is independent of time: each row summarizes one entity over the selected period. The linecharts (Figure 2) can be time-dependent. Both views cover the full data period by default.




Figure 1: The discovery table, showing the countries tab. Each row summarizes one entity over the selected period, and operators can sort by any combination of columns or select up to ten rows to plot in the linecharts.

The discovery table has four tabs, one per dimension: countries, IP addresses, ASN organizations, and usernames. In Figure 1 the countries tab is active. Most columns are shared across tabs: total attacks, average daily attacks, persistence (the share of days on which the entity was active), the largest absolute and percentage day-over-day change, attacks in the most recent seven days, first- and last-seen dates, and the maximum daily count. Some columns apply only to particular tabs; for instance, the countries tab offers a Unique ASNs column that the ASN tab lacks. Operators choose which columns to display through the Columns menu. Each entity also carries a badge summarizing its dominant behavior, namely Scan, Spray, Botnet or Brute [describe how badges are assigned]. A search box filters rows by name, and the counter above the table reports the number of entities in the current tab (197 countries in Figure 1).

Unlike the linecharts, which show only the ten most active entities per dimension, the table lists every entity. Operators can select up to ten rows with the checkboxes and use Analyze Selected to plot those entities in the linecharts, and they can return to the linecharts at any time. This link lets operators move from a precise, ranked lookup to the temporal behavior of the same entities, fulfilling R5.

A.2 Linecharts

The linechart view (Figure 2) stacks five charts on a shared daily time axis. The first, Total Attacks Over Time, shows the aggregate volume of attempts; the remaining four show the ten most active countries, ASN organizations, IP addresses, and usernames. Each legend entry lists the entity’s total for the selected period.

Interaction with the linecharts follows Shneiderman’s mantra [56] (R6), and the charts are coordinated so that an action in any one of them updates all of them.

Overview first. On opening, the charts cover all 69 days of data, with each per-entity chart showing the ten most active entities in its dimension.




Figure 2: The linechart view. The aggregate chart at the top supports brushing to zoom; the four charts below show the ten most active countries, ASN organizations, IP addresses, and usernames.

Zoom. Every chart supports brushing. When the operator brushes a time range on any chart, the range is applied to all five charts and recorded in the Active Filters bar, so the tempo-

Rank	Username	Total Attacks	Avg Daily	Persistence	Max Absolute Δ [1]	Max % Δ [2]	Recent (7d)	First Seen	Last Seen	Max Daily
1	karaf	77.6K	5.2K	21.70% (15d)	32.4K	360.52K%	120	Nov 2, 2022	Jan 8, 2023	44.8K
2	ossuser	82.6K	2.6K	46.40% (32d)	38.4K	276.15K%	690	Nov 3, 2022	Jan 8, 2023	46.0K
3	snow	248.2K	5.2K	69.60% (48d)	22.0K	168.87K%	41.8K	Nov 1, 2022	Jan 8, 2023	35.2K
4	nmrsu	89.4K	1.8K	72.50% (50d)	35.0K	145.91K%	2.9K	Nov 5, 2022	Jan 8, 2023	43.8K
5	lenovo	25.3K	476.51	76.80% (53d)	8.5K	283.19K%	5.7K	Nov 1, 2022	Jan 8, 2023	8.5K
6	fb	9.7K	225.67	62.30% (43d)	4.7K	470.10K%	403	Nov 5, 2022	Jan 8, 2023	4.7K
7	inspur	40.4K	684.71	85.50% (59d)	5.3K	263.05K%	11.2K	Nov 1, 2022	Jan 8, 2023	5.3K
8	radware	12.3K	686.11	26.10% (18d)	4.8K	241.25K%	1.9K	Nov 1, 2022	Jan 8, 2023	4.8K

Figure 3: The usernames tab sorted by Max Absolute Δ (primary key) and Max % Δ (secondary key), surfacing bursty, usernames that may never appear among the top ten by volume.

ral context stays the same whichever dimension the operator is examining. Brushing is disabled once the range is two days or fewer, since a line needs more points to show a trend.

Filter. Left-clicking a legend entry selects that entity and filters every chart to the attacks associated with it. Because each attack attempt links a username, an IP address, an AS, and a country, the selection propagates along these relationships:

- Selecting a country restricts the AS, IP, and username charts to the entities observed in attacks from that country.
- Selecting an AS shows the country it originates from, along with its IP addresses and the usernames they try.
- Selecting an IP address shows its country and AS, along with the usernames it tries.
- Selecting a username shows the countries, ASes, and IP addresses from which it was tried.

Right-clicking a legend entry crosses the entity out instead, removing it from its chart and rescaling the y-axis so that smaller series are no longer flattened against the axis by a dominant one.

Details-on-demand. Hovering over a data point shows a tooltip with that entity’s attack count for the day, hovering over a legend entry highlights the entity’s line, and the legend reports each entity’s total for the selected period.

Linked filtering turns the linecharts from four separate rankings into a means of tracing an entity across dimensions. An operator investigating the spike on 30 December, for example, can select `root` in the username chart; the IP, AS, and country charts then show which addresses, networks, and countries tried `root`, and brushing around 30 December narrows the view to the spike itself. Conversely, selecting an AS such as DigitalOcean reveals which usernames its hosts try, indicating whether traffic from a hosting provider reflects one campaign or many. Brushing on every chart, linked filtering, and on-demand details together fulfill R6.

The linecharts answer questions that the discovery table cannot, such as “Do the countries in Figure 1 share the same peaks and valleys?” or “Does one country begin attacking earlier than another?” In Figure 2, the aggregate chart shows a sharp spike on 30 December, and most of the top ten countries peak on the same day, indicating a coordinated, multi-country event rather than activity from a single source. The

charts also reveal differences in timing. China’s largest peak falls in late November, weeks before the United States’ activity rises, and the IP chart shows individual addresses that each spike for a day or two and then fall silent. Neither pattern can be recovered from the table’s per-entity totals. Together, Figures 1 and 2 show how the two views complement each other, fulfilling R5.




Figure 4: The aggregate, country, and ASN charts at full size. Each per-entity chart shows at most ten series, and the aggregate chart serves as a reference for judging each entity’s share of total attacks.

A.3 Sorting and Ranking

Figure 3 shows the usernames tab sorted on two columns. Operators add sort keys by holding Ctrl/Cmd while clicking column headers, and the bracketed indices in the headers show when the column was clicked.

Sorting matters because the data is heavy-tailed. The username linechart in Figure 2 is dominated by `root`, which accounts for tens of millions of attempts and flattens every other line. Ranking by day-over-day change instead surfaces usernames such as `karaf`, `ossuser`, `lenovo`, `inspur`, and `radware`, several of which are default accounts for specific software or hardware products (`karaf`, for example, is the




Figure 5: Similarity search for the username `james`. The cluster summary, shared infrastructure panel, and “Why similar?” labels explain why the results were grouped with the query.




Figure 6: Similarity search for the IP address `61.177.173.46`. The nearest neighbors are adjacent addresses on the same network

default account of Apache Karaf). Their totals are far below the top ten, but their activity is concentrated: `karaf` received 77.6K attempts in total, 44.8K of them on a single day, and was active on only 15 of the 69 days. Bursts of product-specific credentials like these suggest campaigns targeting particular devices, yet they are invisible in a volume-ranked chart. Figure 3 thus shows how sorting and ranking over the full entity list fulfills R2.

A.4 Visual Clarity

Figure 4 shows the upper portion of the linechart view at full size. Each per-entity chart is limited to ten series, because the accuracy with which people read multi-series line charts declines as the number of series grows [55]. Each chart has its own vertical scale, so that entities are legible even in dimensions with lower totals.

The aggregate chart serves as a reference for the charts below it (R3). On 30 December the aggregate chart peaks at roughly 9.5M attempts, while the username chart in Figure 2 shows `root` at roughly 5M on the same day, so a single username accounts for about half of the spike. Comparisons across

per-entity charts are similarly direct. The United States and DigitalOcean both rise sharply on 9 December and remain elevated for the rest of the period, and a spike from Singapore on 24 November coincides with a spike from Tencent in the ASN chart, consistent with attacks launched from cloud infrastructure hosted in Singapore. Because each chart shows at most ten lines, such correspondences can be spotted at a glance.

B Non-Visual Design

B.1 Username Similarity

Similarity search lets operators start from one attacker and retrieve the attackers that behave most like it. For usernames, the unit of analysis is the population of sources that tries a given username. For each username we compute the breadth of that population (the number of distinct IP addresses, ASes, and countries that try it), its volume, its persistence, and how concentrated its traffic is on its single most active source. We do not reuse the seventeen-dimensional profile we apply to IP addresses (Section B.2), because a diagnostic analysis showed several of those features to be uninformative for usernames. Day-to-day stability is the clearest example: most IP addresses




Figure 7: The empty tag repository for the username dimension.




Figure 8: A newly created tag, listed with its description, author, and application counts.

are active on many days, so their stability is well defined, but about half of the usernames appear on only a single day, where stability is meaningless. We cluster the username profiles with k -means, selecting $k = 9$. Each cluster is labeled automatically from its centroid. To answer a query, the system retrieves the query’s cluster and ranks the other members by Euclidean distance to the query, reported as a similarity score.

In Figure 5, an operator searches for the username `james`. The system places it in Cluster 3, labeled High-Volume Persistent Botnet-Scale, which contains 2,624 usernames; the cluster summary reports that a typical member receives about 31.8K attempts from roughly 4,000 IP addresses and is active on 72% of days. The results table lists the 20 most similar usernames, each with its total attacks, top source country, top three source IPs, and a set of “Why similar?” labels, such as a shared lead source country or comparable volume, that explain the match.

Above the results, a shared infrastructure panel lists IP addresses that appear among the top sources of more than one result, and these addresses are highlighted in the table. In Figure 5, ten addresses are shared. Two of them, 152.89.196.123 and 152.89.196.220, lie in the same /24 network and are top sources for `manager`, `daniel`, `oscar`, `sftp`, `admin1`, and several other usernames. The panel turns behavioral resemblance into infrastructural evidence: usernames tried by the same hosts are likely drawn from the same credential list and belong to the same campaign.

B.2 IP Address Similarity

We characterize each IP address by a seventeen-dimensional behavioral profile drawn entirely from its observed activity during the data period. The features fall into four groups: volume (log-scaled totals, daily means, and peaks); temporal shape (the fraction of days on which the address was active, its burst intensity, its largest day-over-day swings, a fitted trend slope, its activity span, and the share of its traffic in the final week); targeting behavior (the number of distinct usernames it tries, how stable that set is from one day to the next, its rota-

tion rate, and the share of attempts taken by its most-tried username); and two flags indicating whether the address belongs to a cloud provider and whether it originates from a major source country. We cluster these profiles with k -means, label each cluster from its centroid, and answer queries in the same way as for usernames.

In Figure 6, an operator searches for 61.177.173.46, which falls in Cluster 5, labeled Low-Volume Persistent Exploratory, containing 4,405 addresses. The most similar results are 61.177.173.53, 61.177.173.36, 61.177.173.47, and 61.177.172.124, with similarity scores between 91.2% and 80.4%, all belonging to CHINANET-BACKBONE. No feature in the profile encodes an address’s network prefix or AS, yet the nearest neighbors of the query are adjacent addresses on the same network. Behavioral features alone therefore recover infrastructure relationships, suggesting that these hosts are operated together.

Figures 5 and 6 show how SSHVisual lets operators assemble sets of similar attackers along both the username and IP dimensions and see why those attackers were grouped, fulfilling R1.

B.3 Shared Tagging

Figures 7–10 show how an operator records a finding so that other operators can see work that has already been done. Tags are organized in a tag repository for each dimension. The repository displays the current operator’s name, and every tag and tag application records its author.

In Figure 7 the repository is empty. To create a tag, the operator gives it a name, describes what it means, and picks a color; the interface prompts for the description with the reminder that it is what saves the next operator time. In Figure 8 the operator has created a tag, “similarity to james,” for the username cluster found in Figure 5. The repository now lists the tag with its description, its author, and two counts: applications in the current dimension and applications across all dimensions, both zero so far.

7 selected		similarity to james	recurring IP addresses				Apply tag	
☐	Rank	Username	Similarity	Total Attacks	Top Source	Top IPs	Why Similar?	Tags
☐	1	jack	93.6%	71.4K 9.2K IPs	United States	46.101.154.177 (2.5%) 144.24.238.59 (0.5%) 117.50.13.29 (0.3%)	Same lead source (United States) Comparable volume Botnet-scale sourcing (~9,764 IPs)	—
☒	2	abc	93.1%	70.1K 9.2K IPs	United States	68.183.233.225 (2.3%) 68.183.223.15 (1.1%) 45.93.201.82 (1%)	Same lead source (United States) Comparable volume Botnet-scale sourcing (~9,312 IPs)	—
☒	3	jason	91.3%	45.1K 8.2K IPs	United States	202.53.94.150 (2.8%) 209.146.20.162 (1.4%) 85.208.240.169 (0.5%)	Same lead source (United States) Comparable volume Botnet-scale sourcing (~6,170 IPs)	—
☒	4	manager	90.5%	87.4K 9.6K IPs	United States	152.89.196.123 (4.7%) 152.89.196.220 (4.5%) 124.43.129.91 (2.2%)	Same lead source (United States) Comparable volume Botnet-scale sourcing (~9,647 IPs)	—
☒	5	info	90.4%	94.2K 10.9K IPs	United States	143.244.158.100 (1.6%) 144.91.193.44 (0.6%) 89.91.164.55 (0.6%)	Same lead source (United States) Comparable volume Botnet-scale sourcing (~10,945 IPs)	—
☒	6	daniel	90.2%	101.0K 10.9K IPs	United States	152.89.196.123 (3.3%) 152.89.196.220 (3.2%) 141.98.10.63 (0.7%)	Same lead source (United States) Comparable volume Botnet-scale sourcing (~10,926 IPs)	—
☒	7	bin	89.8%	58.3K 8.2K IPs	India	124.43.129.91 (2.6%) 202.53.94.150 (2.6%) 45.93.201.82 (2.1%)	Comparable volume Botnet-scale sourcing (~8,223 IPs) Visible attacker set	—
☒	8	oscar	89.5%	61.8K 7.0K IPs	India	152.89.196.123 (5.3%) 152.89.196.220 (5.3%) 209.146.20.162 (1%)	Comparable volume Botnet-scale sourcing (~7,010 IPs) Persistent (97% of days)	—

Figure 9: Applying the tag to seven usernames from the similarity results, with a note explaining the application.

Tag Repository 1 tag, 7 applications		Operator: Shreeman Change Hide
Tag name	What does it mean? (this is what saves the next operator time)	Colour
e.g. de-firstnames	German first-name wordlist, 3 shared IPs in 109,239,58-60	Create tag
similarity to james seen high similarity for james, which is pretty unusual		7 here / 7 total by Shreeman Delete
Tagged usernames similarity to james Refresh		
abc	similarity to james recurring IP addresses	Shreeman 2026-09-04 05:30 Untag
jason	similarity to james recurring IP addresses	Shreeman 2026-09-04 05:30 Untag
manager	similarity to james recurring IP addresses	Shreeman 2026-09-04 05:30 Untag
info	similarity to james recurring IP addresses	Shreeman 2026-09-04 05:30 Untag
daniel	similarity to james recurring IP addresses	Shreeman 2026-09-04 05:30 Untag
bin	similarity to james recurring IP addresses	Shreeman 2026-09-04 05:30 Untag
oscar	similarity to james recurring IP addresses	Shreeman 2026-09-04 05:30 Untag

Figure 10: The tag repository after the tag is applied, listing each tagged username with its note, author, and timestamp.

Tags are applied from the similarity results. In Figure 9 the operator selects seven usernames (abc, jason, manager, info, daniel, bin, and oscar), chooses the new tag, and adds a note explaining this particular application: “recurring IP addresses.” The selection reflects the evidence from the shared infrastructure panel. All seven usernames have highlighted, shared IP addresses among their top sources, whereas jack, despite the highest similarity score, has none and is left untagged. In Figure 10 the repository reports seven applications. Filtering the list by tag shows each tagged username alongside its note, the operator who applied it, and a timestamp, and any application can be removed with Untag. Tagged usernames also appear in the Tags column of subsequent similarity searches, so an operator who later retrieves the same usernames sees immediately that they have already been investigated.

V Implementation – how we made the tool. this section may not exist by the end

VI Evaluation and Discussion

A Case Studies

B Discussion

VII Summary

A Limitations

B Future Work

VIII Acknowledgments

References

- [1] T. Ylonen, “Ssh—secure login connections over the internet,” in *Proceedings of the 6th USENIX Security Symposium*, vol. 37, pp. 40–52, 1996.
- [2] S. K. Singh, S. Gautam, C. Cartier, S. Patil, and R. Ricci, “Where the wild things are: Brute-Force SSH attacks in

- the wild and how to stop them,” in *21st USENIX Symposium on Networked Systems Design and Implementation (NSDI 24)*, (Santa Clara, CA), pp. 1731–1750, USENIX Association, Apr. 2024.
- [3] M. Antonakakis, T. April, M. Bailey, M. Bernhard, E. Bursztein, J. Cochran, Z. Durumeric, J. A. Halderman, L. Invernizzi, M. Kallitsis, *et al.*, “Understanding the mirai botnet,” in *26th USENIX security symposium (USENIX Security 17)*, pp. 1093–1110, 2017.
 - [4] C. Kolias, G. Kambourakis, A. Stavrou, and J. Voas, “Ddos in the iot: Mirai and other botnets,” *computer*, vol. 50, no. 7, pp. 80–84, 2017.
 - [5] F. Dang, Z. Li, Y. Liu, E. Zhai, Q. A. Chen, T. Xu, Y. Chen, and J. Yang, “Understanding fileless attacks on linux-based iot devices with honeycloud,” in *Proceedings of the 17th Annual International Conference on Mobile Systems, Applications, and Services*, pp. 482–493, 2019.
 - [6] S. Pastrana and G. Suarez-Tangil, “A first look at the crypto-mining malware ecosystem: A decade of unrestricted wealth,” in *Proceedings of the Internet Measurement Conference*, pp. 73–86, 2019.
 - [7] V. T. Guimarães, C. M. D. S. Freitas, R. Sadre, L. M. R. Tarouco, and L. Z. Granville, “A survey on information visualization for network and service management,” *IEEE Communications Surveys & Tutorials*, vol. 18, no. 1, pp. 285–323, 2016.
 - [8] H. Shiravi, A. Shiravi, and A. A. Ghorbani, “A survey of visualization systems for network security,” *IEEE Transactions on Visualization and Computer Graphics*, vol. 18, no. 8, pp. 1313–1329, 2012.
 - [9] T. Zhang, X. Wang, Z. Li, F. Guo, Y. Ma, and W. Chen, “A survey of network anomaly visualization,” *Science China Information Sciences*, vol. 60, no. 12, p. 121101, 2017.
 - [10] J. Raynor, T. Crnovrsanin, S. Di Bartolomeo, L. South, D. Saffo, and C. Dunne, “The state of the art in bgp visualization tools: A mapping of visualization techniques to cyberattack types,” *IEEE Transactions on Visualization and Computer Graphics*, vol. 29, no. 1, pp. 1059–1069, 2022.
 - [11] C. Wu, S. Sheng, and X. Dong, “Research on visualization systems for ddos attack detection,” in *2018 IEEE International Conference on Systems, Man, and Cybernetics (SMC)*, pp. 2986–2991, IEEE, 2018.
 - [12] C. Gates and C. Taylor, “Challenging the anomaly detection paradigm: A provocative discussion,” in *Proceedings of the 2006 workshop on New security paradigms*, pp. 21–29, 2006.
 - [13] F. Mansmann and S. Mansmann, “Interactive exploration of data traffic with hierarchical network maps,” 2006.
 - [14] J. Oberheide, M. Goff, and M. Karir, “Flamingo: Visualizing internet traffic,” in *2006 IEEE/IFIP Network Operations and Management Symposium NOMS 2006*, pp. 150–161, IEEE, 2006.
 - [15] F. Mansmann, F. Fischer, D. A. Keim, and S. C. North, “Visual support for analyzing network traffic and intrusion detection events using treemap and graph representations,” in *Proceedings of the Symposium on Computer Human Interaction for the Management of Information Technology*, pp. 19–28, 2009.
 - [16] T. Taylor, D. Paterson, J. Glanfield, C. Gates, S. Brooks, and J. McHugh, “Flovis: Flow visualization system,” in *2009 Cybersecurity Applications & Technology Conference for Homeland Security*, pp. 186–198, IEEE, 2009.
 - [17] S. Chen, C. Guo, X. Yuan, F. Merkle, H. Schaefer, and T. Ertl, “Oceans: Online collaborative explorative analysis on network security,” in *Proceedings of the Eleventh Workshop on Visualization for Cyber Security*, pp. 1–8, 2014.
 - [18] T. Tremel, J. Kogel, F. Jauernig, S. Meier, D. Thom, F. Becker, C. Muller, and S. Koch, “Vitalflow: Visual interactive traffic analysis with netflow,” in *NOMS 2022-2022 IEEE/IFIP Network Operations and Management Symposium*, pp. 1–6, IEEE, 2022.
 - [19] L. Xiao, J. Gerth, and P. Hanrahan, “Enhancing visual analysis of network traffic using a knowledge representation,” in *2006 IEEE symposium on visual analytics science and technology*, pp. 107–114, IEEE, 2006.
 - [20] J. R. Goodall and M. Sowul, “Viassist: Visual analytics for cyber defense,” in *2009 IEEE conference on technologies for homeland security*, pp. 143–150, IEEE, 2009.
 - [21] Y. Livnat, J. Agutter, S. Moon, and S. Foresti, “Visual correlation for situational awareness,” in *IEEE Symposium on Information Visualization, 2005. INFOVIS 2005.*, pp. 95–102, IEEE, 2005.
 - [22] A. King, B. Huffaker, A. Dainotti, and K. Claffy, “A coordinated view of the temporal evolution of large-scale internet events,” *Computing*, vol. 96, no. 1, pp. 53–65, 2014.
 - [23] D. H. Jeong, B.-K. Jeong, and S.-Y. Ji, “Multi-resolution analysis with visualization to determine network attack patterns,” *Applied Sciences*, vol. 13, no. 6, p. 3792, 2023.
 - [24] D. H. Jeong, J.-H. Cho, F. Chen, L. Kaplan, A. Jøsang, and S.-Y. Ji, “Interactive web-based visual analysis on network traffic data,” *Information*, vol. 14, no. 1, p. 16, 2022.

- [25] D. Ceneda, M. Di Bartolomeo, V. Di Donato, M. Patrignani, M. Pizzonia, and M. Rimondini, "Routingwatch: Visual exploration and analysis of routing events," in *NOMS 2016-2016 IEEE/IFIP Network Operations and Management Symposium*, pp. 591–597, IEEE, 2016.
- [26] J. C. Roberts, "State of the art: Coordinated & multiple views in exploratory visualization," in *Fifth international conference on coordinated and multiple views in exploratory visualization (CMV 2007)*, pp. 61–71, IEEE, 2007.
- [27] J. S. Yi, Y.-a. Kang, J. T. Stasko, and J. A. Jacko, "Understanding and characterizing insights: how do people gain insights using information visualization?," in *Proceedings of the 2008 Workshop on BEyond time and errors: novel evaluation methods for Information Visualization*, pp. 1–6, 2008.
- [28] H. Choi, H. Lee, and H. Kim, "Fast detection and visualization of network attacks on parallel coordinates," *computers & security*, vol. 28, no. 5, pp. 276–288, 2009.
- [29] C. Humphries, N. Prigent, C. Bidan, and F. Majorczyk, "Elvis: Extensible log visualization," in *Proceedings of the Tenth Workshop on Visualization for Cyber Security*, pp. 9–16, 2013.
- [30] T. Yu, R. Lippmann, J. Riordan, and S. Boyer, "Ember: a global perspective on extreme malicious behavior," in *Proceedings of the seventh international symposium on visualization for cyber security*, pp. 1–12, 2010.
- [31] R. Djap, C. Lim, K. E. Silaen, and A. Yusuf, "Xb-pot: Revealing honeypot-based attacker's behaviors," in *2021 9th International Conference on Information and Communication Technology (ICoICT)*, pp. 550–555, IEEE, 2021.
- [32] V. Visoottiviseth, U. Jaralrungruj, E. Phoomrungraungsuk, and P. Kultanon, "Distributed honeypot log management and visualization of attacker geographical distribution," in *2011 Eighth International Joint Conference on Computer Science and Software Engineering (JCSSE)*, pp. 23–28, IEEE, 2011.
- [33] M. Baykara and R. Das, "A novel honeypot based security approach for real-time intrusion detection and prevention systems," *Journal of Information Security and Applications*, vol. 41, pp. 103–116, 2018.
- [34] A. Belqruch and A. Maach, "Scada security using ssh honeypot," in *Proceedings of the 2nd International Conference on Networking, Information Systems & Security*, pp. 1–5, 2019.
- [35] I. Koniaris, G. Papadimitriou, and P. Nicopolitidis, "Analysis and visualization of ssh attacks using honeypots," in *Eurocon 2013*, pp. 65–72, IEEE, 2013.
- [36] A. Kyriakou and N. Sklavos, "Container-based honeypot deployment for the analysis of malicious activity," in *2018 Global Information Infrastructure and Networking Symposium (GIIS)*, pp. 1–4, IEEE, 2018.
- [37] I. Koniaris, G. Papadimitriou, P. Nicopolitidis, and M. Obaidat, "Honeypots deployment for the analysis and visualization of malware activity and malicious connections," in *2014 IEEE international conference on communications (ICC)*, pp. 1819–1824, IEEE, 2014.
- [38] D. Fraunholz, M. Zimmermann, S. D. Anton, J. Schneider, and H. D. Schotten, "Distributed and highly-scalable wan network attack sensing and sophisticated analysing framework based on honeypot technology," in *2017 7th International Conference on Cloud Computing, Data Science & Engineering-Confluence*, pp. 416–421, IEEE, 2017.
- [39] J. R. Goodall, E. D. Ragan, C. A. Steed, J. W. Reed, G. D. Richardson, K. M. Huffer, R. A. Bridges, and J. A. Laska, "Situ: Identifying and explaining suspicious behavior in networks," *IEEE transactions on visualization and computer graphics*, vol. 25, no. 1, pp. 204–214, 2018.
- [40] Q. Liao and A. Striegel, "Intelligent network management using graph differential anomaly visualization," in *2012 IEEE Network Operations and Management Symposium*, pp. 1008–1014, IEEE, 2012.
- [41] W. A. Pike, C. Scherrer, and S. Zabriskie, "Putting security in context: Visual correlation of network activity with real-world information," in *VizSEC 2007: Proceedings of the Workshop on Visualization for Computer Security*, pp. 203–220, Springer, 2008.
- [42] L. Hao, C. G. Healey, and S. E. Hutchinson, "Ensemble visualization for cyber situation awareness of network security data," in *2015 IEEE Symposium on Visualization for Cyber Security (VizSec)*, pp. 1–8, IEEE, 2015.
- [43] R. Sommer and V. Paxson, "Outside the closed world: On using machine learning for network intrusion detection," in *2010 IEEE symposium on security and privacy*, pp. 305–316, IEEE, 2010.
- [44] S. Axelsson, "The base-rate fallacy and the difficulty of intrusion detection," *ACM Transactions on Information and System Security (TISSEC)*, vol. 3, no. 3, pp. 186–205, 2000.
- [45] D. M. Best, R. P. Hafen, B. K. Olsen, and W. A. Pike, "Atypical behavior identification in large-scale network traffic," in *2011 IEEE Symposium on Large Data Analysis and Visualization*, pp. 15–22, IEEE, 2011.
- [46] L. Bilge and T. Dumitras, "Before we knew it: an empirical study of zero-day attacks in the real world," in *Proceedings of the 2012 ACM conference on Computer and communications security*, pp. 833–844, 2012.

- [47] C. Sabottke, O. Suci, and T. Dumitras, “Vulnerability disclosure in the age of social media: Exploiting twitter for predicting real-world exploits,” in *24th USENIX security symposium (USENIX security 15)*, pp. 1041–1056, 2015.
- [48] L. Allodi and F. Massacci, “A preliminary analysis of vulnerability scores for attacks in wild: The ekits and sym datasets,” in *Proceedings of the 2012 ACM Workshop on Building analysis datasets and gathering experience returns for security*, pp. 17–24, 2012.
- [49] Y. Xie, F. Yu, K. Achan, E. Gillum, M. Goldszmidt, and T. Wobber, “How dynamic are ip addresses?,” in *Proceedings of the 2007 conference on Applications, technologies, architectures, and protocols for computer communications*, pp. 301–312, 2007.
- [50] B. Stone-Gross, M. Cova, L. Cavallaro, B. Gilbert, M. Szydlowski, R. Kemmerer, C. Kruegel, and G. Vigna, “Your botnet is my botnet: analysis of a botnet takeover,” in *Proceedings of the 16th ACM conference on Computer and communications security*, pp. 635–647, 2009.
- [51] G. C. Moura, C. Ganán, Q. Lone, P. Poursaied, H. Asghari, and M. van Eeten, “How dynamic is the isps address space? towards internet-wide dhcp churn estimation,” in *2015 IFIP Networking Conference (IFIP Networking)*, pp. 1–9, IEEE, 2015.
- [52] J. L. Guerra, E. Veas, and C. A. Catania, “A study on labeling network hostile behavior with intelligent interactive tools,” in *2019 IEEE Symposium on Visualization for Cyber Security (VizSec)*, pp. 1–10, IEEE, 2019.
- [53] J. Owens and J. Matthews, “A study of passwords and methods used in brute-force SSH attacks,” in *USENIX Workshop on Large-Scale Exploits and Emergent Threats (LEET)*, (San Francisco, CA, USA), 2008.
- [54] Y. Wu, P. M. Cao, A. Withers, Z. T. Kalbarczyk, and R. K. Iyer, “Mining threat intelligence from billion-scale SSH brute-force attacks,” in *Workshop on Decentralized IoT Systems and Security (DISS)*, (San Diego, CA, USA), Internet Society, 2020.
- [55] W. Javed, B. McDonnell, and N. Elmqvist, “Graphical perception of multiple time series,” *IEEE Transactions on Visualization and Computer Graphics*, vol. 16, no. 6, pp. 927–934, 2010.
- [56] B. Shneiderman, “The eyes have it: A task by data type taxonomy for information visualizations,” in *Proceedings of the IEEE Symposium on Visual Languages*, pp. 336–343, 1996.